

Technical Appendix for SPRiF: Spectral Projective Reset Integrate-and-Fire Neuron

Anonymous submission

1 Supplementary Guide and Evidence Map

This technical appendix expands the evidence used in the main paper rather than introducing a second line of claims. It provides reproducibility details for the five benchmark results, exact definitions for the mechanism variants and SI-DMS intervention, and the formal SPRiF update plus dynamical-analysis protocols.

Consistency locks. All sections preserve the main-paper contract: the slow spectral state is three-dimensional, the fast discharge state is two-dimensional, and each timestep follows $I_t \rightarrow \mathbf{x}_t \rightarrow \tilde{\mathbf{u}}_t \rightarrow v_t \rightarrow z_t \rightarrow \mathbf{u}_t$. Projective reset acts directly only on the fast discharge state, with $\Delta \mathbf{x}_t^{\text{reset}} = \mathbf{0}$; in recurrent layers, spikes may still affect the next slow-state update through the ordinary $W_{\text{rec}} \mathbf{z}_t$ input pathway. We use the main-paper terms *functional state decomposition*, *slow spectral state*, *fast discharge state*, and *projective reset* without redefining them. Benchmark SPRiF entries are five-seed mean $\pm$ standard deviation, SI-DMS entries are three-seed summaries, and cross-paper comparator entries remain descriptive comparisons.

Evidence map. Each section below expands a specific main-paper element rather than introducing a second line of claims:

- Main-paper Table 1 is expanded by the dataset encodings, architectures, optimization settings, and reporting rules in Sections 2–2.2; the SPRiF entries remain five-seed estimates and published comparator entries remain descriptive context.
- Main-paper Table 2 is expanded by the variant definitions in Section 3.1: A is $\omega = 0$, B is the merged slow/fast state that changes both state role and total state dimension, and C is $\lambda = 0$.
- The SI-DMS results are expanded by the task construction, shared controls, delay-by- K coverage, and integrity audit in Sections 3.2–3.4, with training intervention counts $K \in \{0, 1, 2, 4\}$ and evaluation intervention counts $K \in \{0, 1, 2, 4, 8\}$.

This is an anonymized submission for review purposes only. Distribution, citation, or public sharing of this manuscript is strictly prohibited. Copyright and publication details will appear in the final version if accepted.

- The S-MNIST loss-landscape diagnostic in Section 2.4 compares one saved SPRiF checkpoint with one saved LIF checkpoint on a shared perturbation-coordinate grid. Because it uses one training seed and one random direction pair, it is treated as a checkpoint-level optimization diagnostic rather than a multi-seed generalization result.
- The mechanism and temporal-kernel figures are expanded by the reset identities and impulse/spectral protocols in Sections 4–4.5.
- A fixed-checkpoint feature-perturbation diagnostic is reported separately in Section 4.6; it records how saved checkpoints respond under feature perturbations without adding a new benchmark, robustness ranking, or biological interpretation for λ .

2 Experimental Protocols and Reporting

The five benchmark entries use task-standard inputs and readouts. S-MNIST and PS-MNIST use MNIST (LeCun et al. 1998) tensors read as 784-step pixel sequences; PS-MNIST applies a seed-generated random pixel permutation fixed within the run. QTDB uses the public QT Database (Laguna et al. 1997) and follows the ASRNN encoding protocol (Yao et al. 2021): two normalized ECG channels are converted by increasing and decreasing level-crossing events into four input event streams, and a time-distributed classifier predicts six waveform classes at each timestep. The released training scripts consume the encoded `QTDB_train.mat` and `QTDB_test.mat` files; this appendix does not claim to provide a raw QTDB-to-`.mat` converter. GSC uses one-second utterances from Google Speech Commands (Warden 2018), represented as 40-bin mel features with first- and second-order deltas stacked over 101 frames, followed by temporal mean pooling over hidden spikes. SHD uses the 700-channel cochlea event sequences from the Spiking Heidelberg Digits dataset (Cramer et al. 2022), with recurrent hidden processing and accumulated readout evidence after warm-up.

2.1 Architectures and Optimization

Table S1 reports the task-specific network and optimization settings for the submitted SPRiF runs. Ranges for ω are expressed as multiples of π .

For S-MNIST and PS-MNIST, truncated BPTT uses a chunk length of 262 steps. The GSC network uses temporal

Item	S-MNIST	PS-MNIST	QTDB	GSC	SHD
Input dim	1	1	4	120	700
Output dim	10	10	6	12	20
Layer	2 Rec (64, 256)	2 Rec (64, 256)	1 Rec (36)	1 Rec (300)	1 Rec (64) + readout
Epochs	150	150	250	150	100
Optimizer	Adam	Adam	Adam	AdamW	Adam
Batch	512	512	128	128	100
LR	1e-2	1e-2	1e-2	5e-3	1e-2
Loss	CE	CE	CE	CE	CE
LR decay	0.1 / 50 ep	0.1 / 50 ep	0.65 / 100 ep	0.5 / 30 ep	0.5 / 20 ep
LR decay type	StepLR	StepLR	StepLR	StepLR	StepLR
Threshold	1.0	1.0	0.6	0.8	1.0
τ_α (steps)	[2, 60]	[2, 60]	[20, 120]	[10, 80]	[10, 80]
τ_ρ (steps)	[2, 40]	[2, 40]	[4, 30]	[4, 30]	[4, 30]
τ_η (steps)	[0.8, 10]	[0.8, 10]	[0.8, 8]	[0.8, 8]	[0.8, 8]
ω ($\times \pi$)	[0.04, 0.50]	[0.04, 0.50]	[0.02, 0.20]	[0.04, 0.40]	[0.04, 0.40]

Table S1: Architecture and optimization settings for the SPRiF benchmark entries. “Rec” denotes a recurrent SPRiF layer, LR decay denotes the StepLR multiplier and interval, and CE denotes cross-entropy.

mean pooling over hidden spikes before the linear classifier. The QTDB network optimizes the time-distributed sequence loss. SHD accumulates readout evidence after the warm-up period. All models use the multi-Gaussian surrogate gradient described in Section 4.3.

2.2 Seeds, Model Selection, and Reporting

All five benchmarks use the seed set $\{0, 1, 2, 42, 1111\}$. For each benchmark and seed, the checkpoint is selected according to validation accuracy, and Table 1 reports the mean and standard deviation of the resulting evaluation accuracies over the five independently seeded runs. Within a task, all runs use the same task-specific preprocessing, data splits, and architecture settings in Table S1; only the seeded training realization changes. Published comparator entries provide cross-paper context rather than matched controls. SI-DMS is a separate controlled mechanism experiment using seeds $\{1, 2, 3\}$ and is reported as a three-seed result.

2.3 Compute Environment

Experiments were run on one NVIDIA GeForce RTX 2080 Ti GPU with 11 GB of memory and 12 vCPUs of an Intel Xeon Platinum 8255C CPU at 2.50 GHz. The software environment used Ubuntu 22.04, Python 3.12, PyTorch 2.3.0, and CUDA 12.1. Training times range from minutes for QTDB to several hours for PS-MNIST and GSC, depending on sequence length and network width. No distributed training or specialized neuromorphic hardware was used.

2.4 S-MNIST Loss-Landscape Diagnostic

We probe local parameter sensitivity around trained S-MNIST solutions using a two-dimensional random-plane visualization inspired by filter-normalized loss landscapes (Li et al. 2018). The SPRiF and LIF models use the same two recurrent layer widths, [64, 256], temporal-mean readout, normalized 784-step S-MNIST inputs, optimizer, learning-rate schedule, and seed. Their neuron dynamics differ, so the

saved models contain 92,810 and 88,970 parameters, respectively; this is therefore an architecture-width-matched rather than parameter-count-matched comparison. The checkpoint metadata record test accuracies of 98.88% for SPRiF and 95.19% for LIF. These values identify the analyzed single-seed checkpoints and do not replace the five-seed benchmark summaries in the main paper.

For each checkpoint θ^* , we generate two Gaussian random directions $\mathbf{d}_1$ and $\mathbf{d}_2$ using seeds 123 and 124. For every multidimensional parameter tensor W , its corresponding direction tensor D is rescaled as $D \leftarrow D \|W\|_F / (\|D\|_F + 10^{-10})$; one-dimensional parameters, including biases and neuron-wise dynamics parameters, are held fixed. We then evaluate

$$\mathcal{L}(a, b) = \mathcal{L}(\theta^* + a\mathbf{d}_1 + b\mathbf{d}_2), \quad (a, b) \in [-1, 1]^2, \quad (1)$$

on a 25×25 grid. Cross-entropy is estimated on the first four non-shuffled test mini-batches (2,048 images), with the 784-step sequence processed in chunks of 262 steps and the chunk-level logits averaged. Both models use the same coordinate range, grid resolution, and evaluation samples; the directions are normalized separately against each model’s checkpoint tensors.

Figure S1 shows a broader low-loss neighborhood for the analyzed SPRiF checkpoint. At the unperturbed center, the estimated cross-entropy is 0.0575 for SPRiF and 0.2098 for LIF. Over the 625 sampled coordinates, the SPRiF slice has lower loss at 85.1% of points, with mean losses of 10.69 versus 15.78 and maxima of 36.46 versus 50.85. Thus, under these norm-matched perturbations, this SPRiF solution is less sensitive along the sampled plane. The result is deliberately scoped to one checkpoint and one direction pair: it does not characterize the full high-dimensional landscape, isolate which SPRiF component creates the geometry, or by itself establish a causal link between flatness and generalization.




Figure S1: S-MNIST loss-landscape diagnostic for SPRiF (top row) and LIF (bottom row). The 2D contours and 3D surfaces show cross-entropy over the same 25×25 perturbation-coordinate grid, with a shared logarithmic color scale and common vertical range. Each model uses its own checkpoint-normalized random directions, while evaluation samples and grid coordinates are matched.

3 Mechanism Variants and SI-DMS Evaluation

3.1 Mechanism Variant Definitions

Table 2 in the main paper reports the ablation outcomes. Table S2 defines the exact intervention behind each column without repeating the main-paper numbers.

3.2 SI-DMS Protocol

The Spike-Intervention Delayed Match-to-Sample (SI-DMS) task tests whether a network retains a cue through a delay while selected hidden units are forced to emit real spikes. Table S3 gives the protocol and matched-control configuration.

3.3 Condition-Level Grids

The condition-level SI-DMS result artifact underlying the main-paper mechanism figure covers all six models, three seeds, five evaluation delays, and five intervention levels. The main-paper SI-DMS table summarizes only clean ($K = 0$) and maximum-stress ($K = 8$) conditions; the full evaluated grid is scoped here by its delay and intervention coverage, while the non-duplicated visualization remains in the main paper. This appendix entry therefore scopes the already-evaluated grid rather than contributing a new experiment, statistic, or robustness claim.

3.4 Model Controls and Integrity Audit

At an intervention timestep, the selected membrane coordinate is placed above threshold before spike computation. Under the hard threshold rule in Eq. (17), this yields an actual spike and then applies the model’s ordinary reset rule. The intervention therefore does not bypass reset and does not add a separate post-spike update.

For SPRiF, let the reset direction be $\mathbf{r} = [1, \lambda_j]^\top$, so the fast-state reset identity is

$$\mathbf{u}_t - (\tilde{\mathbf{u}}_t - z_t \theta \mathbf{r}) = \mathbf{0}, \quad (2)$$

while the slow state has no reset update,

$$\Delta \mathbf{x}_t^{\text{reset}} = \mathbf{0}. \quad (3)$$

These are structural identities of the stated update. They formalize the SI-DMS distinction: forced spikes can alter the fast discharge state through real spike-and-reset events, but they do not directly modify the slow spectral state at that timestep.

4 Formal SPRiF Specification and Analysis Protocols

4.1 Slow Spectral State

The slow spectral state $\mathbf{x}_t \in \mathbb{R}^3$ evolves as

$$\mathbf{x}_t = A \mathbf{x}_{t-1} + B I_t, \quad (4)$$

where

$$A = \begin{bmatrix} \alpha & 0 & 0 \\ 0 & \rho \cos \omega & -\rho \sin \omega \\ 0 & \rho \sin \omega & \rho \cos \omega \end{bmatrix}, \quad B = \begin{bmatrix} 1 - \alpha \\ 1 - \rho \\ 0 \end{bmatrix}. \quad (5)$$

The first coordinate is a real exponential trace,

$$x_t^0 = \alpha x_{t-1}^0 + (1 - \alpha) I_t, \quad (6)$$

and the second and third coordinates form a damped rotational pair driven through the first oscillatory coordinate:

$$x_t^1 = \rho \cos \omega x_{t-1}^1 - \rho \sin \omega x_{t-1}^2 + (1 - \rho) I_t, \quad (7)$$

$$x_t^2 = \rho \sin \omega x_{t-1}^1 + \rho \cos \omega x_{t-1}^2. \quad (8)$$

The eigenvalues of A are α and $\rho e^{\pm i\omega}$. We parameterize

$$\alpha = \sigma(\alpha_{\text{raw}}), \quad \rho = \sigma(\rho_{\text{raw}}), \quad \omega = \pi \sigma(\omega_{\text{raw}}), \quad (9)$$

so $\alpha, \rho \in (0, 1)$ and $\omega \in (0, \pi)$, giving spectral radius $\max(\alpha, \rho) < 1$. These sigmoid parameterizations keep the slow spectral state strictly stable under the learned raw variables.

Under a constant input $I_t = c$, the real mode converges to c . The rotational pair converges to the unique fixed point

$$\mathbf{x}_*^{\text{osc}} = (I - \rho R_\omega)^{-1} \begin{bmatrix} (1 - \rho)c \\ 0 \end{bmatrix}, \quad (10)$$

$$R_\omega = \begin{bmatrix} \cos \omega & -\sin \omega \\ \sin \omega & \cos \omega \end{bmatrix}. \quad (11)$$

Thus the first oscillatory coordinate is not generally equal to c when $\omega \neq 0$.

For a unit input impulse at $t = 0$ with $\mathbf{x}_{-1} = \mathbf{0}$, the slow-state impulse response isolates the two retention modes separately:

$$x_k^0 = (1 - \alpha) \alpha^k, \quad (12)$$

$$x_k^1 = (1 - \rho) \rho^k \cos(k\omega), \quad (13)$$

$$x_k^2 = (1 - \rho) \rho^k \sin(k\omega), \quad (14)$$

for $k \geq 0$. The real mode provides monotone retention, whereas the rotational pair provides damped phase dynamics.

4.2 Fast Discharge State and Projective Reset

The fast discharge state $\mathbf{u}_t \in \mathbb{R}^2$ receives a projection of the slow spectral state through $G \in \mathbb{R}^{2 \times 3}$:

$$\tilde{\mathbf{u}}_t^0 = \eta_0 u_{t-1}^0 + \kappa u_{t-1}^1 + (G \mathbf{x}_t)^0, \quad (15)$$

$$\tilde{\mathbf{u}}_t^1 = \eta_1 u_{t-1}^1 + (G \mathbf{x}_t)^1, \quad (16)$$

with $\eta_0, \eta_1 \in (0, 1)$ parameterized as sigmoid-transformed raw variables. The membrane potential is $v_t = \tilde{u}_t^0$, and the spike is

$$z_t = H(v_t - \theta), \quad (17)$$

where H is the Heaviside step function.

Within each timestep, computation follows $I_t = W_{\text{in}} s_t + W_{\text{rec}} \mathbf{z}_{t-1} + b \rightarrow \mathbf{x}_t = A \mathbf{x}_{t-1} + B I_t \rightarrow \tilde{\mathbf{u}}_t \rightarrow v_t = \tilde{u}_t^0 \rightarrow z_t = H(v_t - \theta) \rightarrow \mathbf{u}_t = \tilde{\mathbf{u}}_t - z_t \theta \mathbf{r}$. This order matches the main paper and confines the local reset update to the fast discharge state.

Component-wise, projective reset is

$$u_t^0 = \tilde{u}_t^0 - z_t \theta, \quad (18)$$

$$u_t^1 = \tilde{u}_t^1 - z_t \theta \lambda_j. \quad (19)$$

Setting $\lambda_j = 0$ recovers scalar subtraction on the membrane coordinate; learned λ_j allows reset to act along $[1, \lambda_j]^\top$ in the fast state.

Variant	Change relative to SPRiF full	Interpretation lock
$\omega = 0$	Removes rotational coupling while retaining the real decay path and slow-fast state separation.	Tests spectral rotation, not state separation.
Merged	Collapses the slow spectral state and fast discharge state into one three-dimensional state whose first coordinate is both membrane and reset target.	Tests the consequence of collapsing memory and discharge; it is not a pure reset-direction control and does not preserve the full model’s total state dimensionality.
$\lambda = 0$	Fixes the reset direction to $[1, 0]^\top$, so reset subtracts only from the membrane coordinate.	Tests learned projective reset against scalar reset while preserving slow-fast separation.

Table S2: Mechanism variants corresponding to the main-paper ablation table.

Item	Setting
Input channels	30 total: 10 left-cue channels, 10 right-cue channels, and 10 noise channels, encoded as Poisson spike trains.
Pre-period	50 ms baseline period with noise channels active.
First / second cue	100 ms each; each cue is left or right.
Delay (train)	Drawn from $\{200, 400, 800, 1600\}$ ms.
Delay (eval)	Drawn from $\{200, 400, 800, 1600, 2500\}$ ms.
Target	Binary match/non-match label depending on whether the two cues are identical.
Intervention	At K random delay timesteps, the membrane coordinates of 10% of hidden units are placed above threshold before spike computation; the resulting spike follows the model’s ordinary reset rule.
K (train)	$\{0, 1, 2, 4\}$.
K (eval)	$\{0, 1, 2, 4, 8\}$, including $K = 8$ as an unseen stress level.
Shared controls	Within each seed/delay/ K condition, models use matched batches, intervention masks, and evaluation probes.
Hidden units	64.
Architecture	Feedforward hidden layer; no recurrent hidden pathway.
SPRiF variants	<code>sprif_full</code> , <code>sprif_lambda0</code> , and <code>sprif_merged</code> .
Baselines	LIF, ASRNN, and BRNN with the same hidden width, optimizer, batch structure, and intervention masks.

Table S3: SI-DMS task construction, controlled intervention, and matched-control configuration.

Audit item	Required behavior	Appendix status
Coverage	Each model/seed evaluates every delay-by- K cell.	Six models, three seeds, five delays, and five K levels are covered in the saved formal evaluation files.
Forced-spike hit	Every selected unit should spike when $K > 0$.	The saved evaluation metrics record forced-hit rate = 1.0 for all 360 intervention cells with $K > 0$.
Shared controls	Comparisons should not depend on model-specific probes.	Within a seed/delay/ K condition, batches and intervention masks are shared across models.
Reset path	Forced spikes should trigger the same reset rule as ordinary spikes.	Interventions are applied pre-threshold; the resulting spike enters the standard reset update.
SPRiF slow-state insulation	A forced spike must not directly edit the slow spectral state.	Eq. (3) states $\Delta \mathbf{x}_t^{\text{reset}} = \mathbf{0}$; recurrent communication can only enter the next timestep through $W_{\text{rec}} \mathbf{z}_t$ when recurrence is present.

Table S4: SI-DMS integrity checks used to interpret the intervention results.

4.3 Surrogate Gradient

The forward pass uses the hard threshold in Eq. (17). During backpropagation, the threshold derivative is replaced by the multi-Gaussian surrogate (Yao et al. 2021):

$$\frac{\partial z_t}{\partial v_t} \approx \gamma [a_0 \mathcal{N}(v_t - \theta; 0, \sigma_0^2) - a_1 \mathcal{N}(v_t - \theta; \mu_1, \sigma_1^2) - a_1 \mathcal{N}(v_t - \theta; -\mu_1, \sigma_1^2)], \quad (20)$$

where $\sigma_0 = 0.5$, $a_0 = 1.15$, $\mu_1 = 0.5$, $\sigma_1 = 3.0$, $a_1 = 0.15$, and $\gamma = 0.5$. The forward dynamics remain unchanged.

4.4 Initialization, State Size, and Cost

The task-specific initialization ranges are listed in Table S1. For each decay $q \in \{\alpha, \rho, \eta_0, \eta_1\}$, a time constant τ_q in steps is mapped to $q = \exp(-1/\tau_q)$ before conversion to its unconstrained raw parameter. Raw variables are initialized so that the transformed α , ρ , η_i , and ω values fall in the listed ranges; after initialization, optimization proceeds in the unconstrained raw parameter space.

Each SPRiF neuron stores five state values $(x^0, x^1, x^2, u^0, u^1)$ and has 13 neuron-wise dynamics parameters: α_{raw} , ρ_{raw} , ω_{raw} , $\eta_{0,\text{raw}}$, $\eta_{1,\text{raw}}$, κ , the six entries of $G \in \mathbb{R}^{2 \times 3}$, and λ . The weight matrices W_{in} and W_{rec} and the bias b follow the layer architecture.

For a layer of width h , the local state update cost is $O(h)$ because each neuron uses fixed-size slow and fast states. When recurrent connectivity is enabled, the $O(h^2)$ multiplication $W_{\text{rec}} \mathbf{z}_{t-1}$ dominates for both LIF-style recurrent SNNs and SPRiF layers.

4.5 Analysis Protocols

For the impulse-response and temporal-kernel analyses in the main paper, we inject a unit impulse into trained SPRiF neurons and then record 100 steps of intrinsic slow-state evolution without further input. The four displayed SPRiF trajectories use the 10th, 40th, 70th, and 90th percentiles of learned α . For the GSC and QTDB cross-model panels, the 90th-percentile SPRiF real mode is paired with the available trained ASRNN membrane kernel having the closest derived τ_m , and each displayed kernel is normalized by its peak magnitude. Spectra from the real coordinate are aggregated across neurons. Because no subsequent input is applied and the slow state has no direct local reset term, the observed response shapes reflect the learned spectral parameters rather than later post-spike correction.

The learned reset-direction analysis uses the λ values shown in the main-paper mechanism figure. Across QTDB, GSC, and PS-MNIST, the reported λ distributions span both signs, and the main-paper figure reports weak Pearson correlations with firing rate, α , and ω ($|r| \leq 0.191$). This appendix treats those observations as evidence that projective reset directions remain diverse in the fast discharge state; it does not attach a universal biological meaning to either sign of λ .

4.6 Feature-Perturbation Diagnostic

Figure S2 reports a fixed-checkpoint feature-space perturbation diagnostic on GSC and QTDB. Controlled input cor-

ruptions are widely used to probe non-adversarial robustness (Hendrycks and Dietterich 2019), while additive Gaussian jittering is a standard magnitude-domain transformation for time-series inputs (Iwana and Uchida 2021). Motivated by these evaluation practices, we apply additive Gaussian noise, random feature zeroing, and their combination to saved SPRiF and ASRNN checkpoints, with matched perturbation settings within each condition. This diagnostic is not a replacement for the five-seed benchmark table in the main paper and is not used as a statistical robustness ranking; it records how the saved checkpoints respond under the specified feature perturbations.

5 Reproducibility Checklist

- **Data sources:** S-MNIST and PS-MNIST use MNIST; QTDB uses the public QT Database encoded following Yin et al. (Yao et al. 2021); GSC uses Google Speech Commands; SHD uses Spiking Heidelberg Digits. See Section 2.
- **Preprocessing:** Dataset encodings, input dimensions, and readout conventions are summarized in Section 2. QTDB is described as an encoded four-event-stream input; no raw QTDB conversion script is introduced here.
- **Architectures and hyperparameters:** Network widths, epochs, optimizer settings, learning-rate schedules, thresholds, and initialization ranges are listed in Table S1.
- **Seeds and reporting:** Main benchmarks use seeds $\{0, 1, 2, 42, 1111\}$ and report mean $\pm$ standard deviation for SPRiF; checkpoints are selected by validation accuracy. SI-DMS uses seeds $\{1, 2, 3\}$; published comparator entries are descriptive cross-paper references. See Section 2.2.
- **Surrogate gradient and dynamics:** The multi-Gaussian surrogate is specified in Section 4.3; the exact slow and fast updates are specified in Sections 4.1 and 4.2.
- **SI-DMS semantics:** Delay sets, K sets, forced-spike semantics, shared controls, condition-level grids, and integrity checks are documented in Sections 3.2–3.4.
- **Loss-landscape diagnostic:** The checkpoint selection, direction normalization, 25×25 grid, test subset, quantitative summaries, and single-slice interpretation boundaries are documented in Section 2.4.
- **Analysis protocols:** The impulse-response, spectral-summary, learned-reset-direction interpretations, and fixed-checkpoint feature-perturbation diagnostic are scoped in Sections 4.5–4.6.
- **Code availability:** An anonymized code archive accompanying the submission contains the five benchmark implementations, all analysis code used in the paper, and the SI-DMS implementation. The complete code will be released under the MIT License upon publication. No identifying paths or author information are included in the submitted archive.




Figure S2: Feature-space perturbation diagnostic using saved checkpoints. SPRiF and ASRNN are evaluated on GSC and QTDB under additive Gaussian noise, random feature zeroing, and mixed perturbations. The clean values in this diagnostic come from the saved checkpoints used for the perturbation sweep and need not equal the five-seed benchmark summaries in the main paper.

References

- Cramer, B.; Stradmann, Y.; Schemmel, J.; and Zenke, F. 2022. The Heidelberg Spiking Data Sets for the Systematic Evaluation of Spiking Neural Networks. *IEEE Transactions on Neural Networks and Learning Systems*, 33(7): 2744–2757.
- Hendrycks, D.; and Dietterich, T. G. 2019. Benchmarking Neural Network Robustness to Common Corruptions and Perturbations. In *International Conference on Learning Representations*.
- Iwana, B. K.; and Uchida, S. 2021. An Empirical Survey of Data Augmentation for Time Series Classification with Neural Networks. *PLOS ONE*, 16(7): e0254841.
- Laguna, P.; Mark, R. G.; Goldberger, A. L.; and Moody, G. B. 1997. A Database for Evaluation of Algorithms for Measurement of QT and Other Waveform Intervals in the ECG. In *Computers in Cardiology*, volume 24, 673–676.
- LeCun, Y.; Bottou, L.; Bengio, Y.; and Haffner, P. 1998. Gradient-Based Learning Applied to Document Recognition. *Proceedings of the IEEE*, 86(11): 2278–2324.
- Li, H.; Xu, Z.; Taylor, G.; Studer, C.; and Goldstein, T. 2018. Visualizing the Loss Landscape of Neural Nets. In *Advances in Neural Information Processing Systems*, volume 31, 6391–6401.
- Warden, P. 2018. Speech Commands: A Dataset for Limited-Vocabulary Speech Recognition. *CoRR*, abs/1804.03209.
- Yao, Z.; Zeng, R.; Peng, Q.; Dong, J.; Xiao, Z.; and Xu, G. 2021. Accurate and Efficient Time-Domain Classification with Adaptive Spiking Recurrent Neural Networks. *Nature Machine Intelligence*, 3: 912–922.